

Optimal PSO from Commutative Weak Pseudorandom Functions

Abstract

XXX

Keywords: XXX

Contents

1	Preliminaries	1
1.1	MPC in the Semi-honest Model	1
1.2	Private Set Union	1
2	Commutative Pseudorandom Functions	1
2.1	Commutative Weak PRFs from the DDH Assumption	2
3	PSU from Commutative Weak PRFs	2
4	Instantiation Based on the DDH Assumption	4
A	Diffie-Hellman Based PSI	7

1 Preliminaries

1.1 MPC in the Semi-honest Model

We use the standard notion of security in the presence of semi-honest adversaries. Let Π be a protocol for computing the function $f(x_1, x_2)$, where party P_i has input x_i . We define security in the following way. For each party P , let $\text{View}_P(x_1, x_2)$ denote the view of party P during an honest execution of Π on inputs x_1 and x_2 . The view consists of P 's input, random tape, and all messages exchanged as part of the Π protocol.

Definition 1.1. 2-party protocol Π securely realizes f in the presence of semi-honest adversaries if there exists a simulator Sim such that for all inputs x_1, x_2 and all $i \in \{1, 2\}$:

$$\text{Sim}(i, x_i, f(x_1, x_2)) \approx_c \text{View}_{P_i}(x_1, x_2)$$

1.2 Private Set Union

PSU is a special case of secure two-party computation. We call the two parties engaging in PSU the *sender* and the *receiver*. The sender holds a set X of size n_x , and the receiver holds a set Y of size n_y . Both sets consist of σ -bit strings. We always assume the set sizes n_x and n_y are public. The ideal PSU functionality (depicted in Figure 1) computes the union, outputs nothing to the sender, and $X \cup Y$ to the receiver.

Parameters:

- Sender \mathcal{S} , Receiver \mathcal{R}
- Set sizes n_x and n_y

Functionality:

- Wait for input $X = \{x_1, \dots, x_{n_x}\} \subset \{0, 1\}^*$ from the receiver \mathcal{S} .
- Wait for input $Y = \{y_1, \dots, y_{n_y}\} \subset \{0, 1\}^*$ from the sender \mathcal{R} .
- Give output $X \cup Y$ to the receiver \mathcal{R} .

Figure 1: Private Set Union Functionality \mathcal{F}_{psu}

2 Commutative Pseudorandom Functions

Let F be a family of pseudorandom functions from D to R . The standard security requirement for PRFs is pseudorandomness.

Pseudorandomness. Let \mathcal{A} be an adversary against PRFs and define its advantage as:

$$\text{Adv}_{\mathcal{A}}(\lambda) = \Pr \left[\begin{array}{l} pp \leftarrow \text{Setup}(1^\lambda); \\ k \leftarrow \text{KeyGen}(pp); \\ b \leftarrow \{0, 1\}; \\ b' \leftarrow \mathcal{A}^{\mathcal{O}_{\text{ror}}(b, \cdot)}(\lambda); \end{array} \right] - \frac{1}{2},$$

where $\mathcal{O}_{\text{ror}}(0, x) = F_k(x)$, $\mathcal{O}_{\text{ror}}(1, x) = \text{H}(x)$ (here H is chosen uniformly at random from all the functions from D to R ¹). Note that \mathcal{A} can adaptively access the oracle $\mathcal{O}_{\text{ror}}(b, \cdot)$ polynomial many times. We say

¹To efficiently simulate access to a uniformly random function H from D to R , one may think of a process in which the adversary's queries to $\mathcal{O}_{\text{ror}}(1, \cdot)$ are “lazily” answered with independently and randomly chosen elements in R , while keeping track of the answers so that queries made repeatedly are answered consistently.

that F is pseudorandom if for any PPT adversary its advantage function $\text{Adv}_{\mathcal{A}}(\lambda)$ is negligible in λ . We refer to such security as full PRF security.

Sometimes the full PRF security is not needed and it is sufficient if the function cannot be distinguished from a uniform random one when challenged on random inputs. The formalization of such relaxed requirement is *weak pseudorandomness*, which is defined the same way as pseudorandomness except that the inputs of oracle $\mathcal{O}_{\text{ror}}(b, \cdot)$ are uniformly chosen from D by the challenger instead of adversarially chosen by \mathcal{A} . PRFs that satisfy weak pseudorandomness are referred to as *weak PRFs*.

Composable. For a family of keyed function F , F is 2-composable if $R \subseteq D$, namely, for any $k_1, k_2 \in K$, the function $F_{k_1}(F_{k_2}(\cdot))$ is well-defined. From now on, we will assume $R = D$ for simplicity.

Commutative. For a family of composable keyed function, we say it is commutative if:

$$\forall k_1, k_2 \in K, \forall x \in X : F_{k_1}(F_{k_2}(x)) = F_{k_2}(F_{k_1}(x))$$

It is easy to see that the standard pseudorandomness denies commutative property, but weak pseudorandomness and commutative property may co-exist.

2.1 Commutative Weak PRFs from the DDH Assumption

We build a concrete commutative weak PRF from the DDH assumption.

- **Setup**(1^λ): runs $\text{GroupGen}(1^\lambda) \rightarrow (\mathbb{G}, g, p)$, outputs $pp = (\mathbb{G}, g, p)$.
- **KeyGen**(pp): outputs $k \xleftarrow{\mathbb{R}} \mathbb{Z}_p$. Each k defines a function from \mathbb{G} to \mathbb{G} , which takes $x \in \mathbb{G}$ as input and outputs x^k .

It is straightforward to verify that F is commutative. By the random self-reducibility of the DDH assumption, the weak pseudorandomness can be tightly reduced to the DDH assumption. We omit the details for triviality.

3 PSU from Commutative Weak PRFs

Parameters:

- Common input: $F : K \times D \rightarrow D$, hash function $H : \{0, 1\}^* \rightarrow D$.
- Input of sender \mathcal{S} : $X = \{x_1, \dots, x_n\}$.
- Input of receiver \mathcal{R} : $Y = \{y_1, \dots, y_n\}$.

Protocol:

1. \mathcal{R} picks $k_2 \xleftarrow{\mathbb{R}} K$, then sends $\{F_{k_2}(H(y_1)), \dots, F_{k_2}(H(y_n))\}$ to \mathcal{S} .
2. \mathcal{S} picks $k_1 \xleftarrow{\mathbb{R}} K$, computes and sends $\{F_{k_1}(H(x_1)), \dots, F_{k_1}(H(x_n))\}$ to \mathcal{R} ; then computes $\{F_{k_1}(F_{k_2}(H(y_1))), \dots, F_{k_1}(F_{k_2}(H(y_n)))\}$, sends its permutation Γ to \mathcal{R} ;
3. \mathcal{R} computes $\{F_{k_2}(F_{k_1}(H(x_1))), \dots, F_{k_2}(F_{k_1}(H(x_n)))\}$, then sets $v_i = 0$ iff the value is not in Γ .
4. \mathcal{R} with select vector (v_1, \dots, v_n) and \mathcal{S} with input $\{(x_i, \perp)\}_{i \in [n]}$ engage in one-sided OT.
5. \mathcal{R} obtains $X - X \cap Y$, and thus obtains the union $X \cup Y$.

Figure 2: PSU from Commutative weak PRFs

Correctness. The above protocol is correct except the case E that $F_{k_1}(F_{k_2}(\mathbf{H}(x))) = F_{k_1}(F_{k_2}(\mathbf{H}(y)))$ for some $x \neq y$ occurs. We further divide E to E_0 and E_1 . E_0 denotes the case that $\mathbf{H}(x) = \mathbf{H}(y)$. E_1 denotes the case that $\mathbf{H}(x) \neq \mathbf{H}(y)$ but $F_{k_1}(F_{k_2}(\mathbf{H}(x))) = F_{k_1}(F_{k_2}(\mathbf{H}(y)))$, which can further be divided into sub-cases $E_{10} \text{ --- } F_{k_2}(\mathbf{H}(x)) = F_{k_2}(\mathbf{H}(y))$ and $E_{11} \text{ --- } F_{k_2}(\mathbf{H}(x)) \neq F_{k_2}(\mathbf{H}(y))$ but $F_{k_1}(F_{k_2}(\mathbf{H}(x))) = F_{k_1}(F_{k_2}(\mathbf{H}(y)))$. By the collision resistance of \mathbf{H} , we have $\Pr[E_0] = 2^{-\sigma}$. By the weak pseudorandomness of F , we have $\Pr[E_{10}] = \Pr[E_{11}] = 2^{-\ell}$. Therefore, we have $\Pr[E] \leq \Pr[E_0] + \Pr[E_{10}] + \Pr[E_{11}] = 2^{-\sigma} + 2^{-\ell+1}$.

Theorem 3.1. *The above PSU protocol is secure in the semi-honest model assuming \mathbf{H} is a random oracle and F is a family of commutative weak PRFs.*

Proof. We exhibit simulators $\text{Sim}_{\mathcal{R}}$ and $\text{Sim}_{\mathcal{S}}$ for simulating corrupt \mathcal{R} and \mathcal{S} respectively, and argue the indistinguishability of the produced transcript from the real execution. Let $|X \cap Y| = m$.

Corrupt sender: $\text{Sim}_{\mathcal{S}}$ simulates the view of corrupt \mathcal{S} , which consists of \mathcal{S} 's randomness, input, output and received messages.

We argue the output of $\text{Sim}_{\mathcal{S}}$ is indistinguishable from the real execution. For this, we formally show the simulation by proceeding the sequence of hybrid transcripts, where T_0 is the real view of \mathcal{S} , and T_2 is the output of $\text{Sim}_{\mathcal{S}}$.

Hybrid₀: $\text{Sim}_{\mathcal{S}}$ simulates with the knowledge of Y .

- $\text{Sim}_{\mathcal{S}}$ chooses the randomness for \mathcal{R} , i.e., picks $k_2 \xleftarrow{\mathbf{R}} K$.
- RO queries: for random oracle query $\langle z \rangle$, picks $h \xleftarrow{\mathbf{R}} D$ and set $\mathbf{H}(z) := h$.
- Let $h_i = \mathbf{H}(y_i)$ for each $y_i \in Y$. $\text{Sim}_{\mathcal{S}}$ outputs $(F_{k_2}(h_1), \dots, F_{k_2}(h_n))$.

Clearly, $\text{Sim}_{\mathcal{S}}$'s simulation is identical to the real view.



Hybrid₁: $\text{Sim}_{\mathcal{S}}$ simulates without the knowledge of Y , and changes the simulation in the final input.

- $\text{Sim}_{\mathcal{S}}$ outputs (s_1, \dots, s_n) where $s_i \xleftarrow{\mathbf{R}} D$.

We argue that the view in hybrid 0 and hybrid 1 are computationally indistinguishable. More precisely, a PPT adversary \mathcal{A} (with knowledge of X and Y) against commutative weak PRF (with secret key k) is given n tuples (h_i, s_i) where $h_i \xleftarrow{\mathbf{R}} D$, and is asked to distinguish if $s_i = F_k(h_i)$ or s_i are random values. \mathcal{A} implicitly sets \mathcal{R} 's randomness $k_2 := k$.

- RO queries: for random oracle query $\langle z \rangle$ where $z \notin Y$, picks $h \xleftarrow{\mathbf{R}} D$ and set $\mathbf{H}(z) := h$; for random oracle query $\langle y_i \rangle$ where $y_i \in Y$, sets $\mathbf{H}(y_i) := h_i$.
- \mathcal{A} outputs (s_1, \dots, s_n) .

If $s_i = F_k(h_i)$ for $i \in [n]$, then \mathcal{A} 's simulation is identical to hybrid 0. If s_i are random values, then \mathcal{A} 's simulation is identical to hybrid 1.

Corrupt receiver: $\text{Sim}_{\mathcal{R}}$ simulates the view of corrupt \mathcal{R} , which consists of \mathcal{R} 's randomness, input, output and received messages.

We argue the output of $\text{Sim}_{\mathcal{R}}$ is indistinguishable from the real execution. For this, we formally show the simulation by proceeding the sequence of hybrid transcripts, where T_0 is the real view of \mathcal{R} , and T_2 is the output of $\text{Sim}_{\mathcal{R}}$.

Hybrid₀: $\text{Sim}_{\mathcal{R}}$ simulates with the knowledge of X .

- $\text{Sim}_{\mathcal{R}}$ chooses the randomness for \mathcal{S} , i.e., picks $k_1 \xleftarrow{\mathbf{R}} K$.

- RO queries: for random oracle query $\langle z \rangle$, picks $h \xleftarrow{R} D$ and sets $H(z) := h$.
- $\text{Sim}_{\mathcal{R}}$ computes and outputs $\{F_{k_1}(H(x_i))\}_{x_i \in X}$, computes $\{F_{k_1}(F_{k_2}(H(y_i)))\}_{y_i \in Y}$, outputs its permutation Γ .

Clearly, $\text{Sim}_{\mathcal{R}}$'s simulation is identical to the real view.



Hybrid₁: $\text{Sim}_{\mathcal{R}}$ simulates without the knowledge of $X \cap Y$, but with the knowledge of $X - X \cap Y$.

- $\text{Sim}_{\mathcal{R}}$ picks $s_j \xleftarrow{R} D$ for $j \in [n - m]$ (associated with PRF values for elements in $X - X \cap Y$), picks $s_k \xleftarrow{R} D$ for $k \in [m]$ (implicitly associated with PRF values for elements in $X \cap Y$, though the exact elements are unknown), outputs $(\{s_j\}_{j \in [n-m]}, \{s_k\}_{k \in [m]})$ according to the order of the final selection vector; picks $s_\ell \xleftarrow{R} D$ for $\ell \in [n - m]$ (associated with PRF values for elements in $Y - X \cap Y$), outputs the permutation of $(\{F_{k_2}(s_k)\}_{k \in [m]}, \{F_{k_2}(s_\ell)\}_{\ell \in [n-m]})$.

We argue that the view in hybrid 1 and hybrid 2 are computationally indistinguishable. More precisely, a PPT adversary \mathcal{A} (with knowledge of X and Y) against commutative weak PRFs are given $n + m$ tuples (h_i, s_i) where $h_i \xleftarrow{R} D$, and is asked to determine if $s_i = F_k(h_i)$ or random values. \mathcal{A} implicitly sets \mathcal{S} 's randomness $k_1 := k$, picks $k_2 \xleftarrow{R} K$.

- RO queries: for $z \notin X \cup Y$, picks $h \xleftarrow{R} D$ and returns $H(z) := h$; for $z_i \in X \cup Y$, returns $H(z_i) = h_i$.
- \mathcal{A} prepares s_j for $z_i \in X - X \cap Y$, s_k for $z_k \in X \cap Y$, outputs $(\{s_j\}_{j \in [n-m]}, \{s_k\}_{k \in [m]})$ according to the order of the final selection vector; prepares $s_\ell \xleftarrow{R} D$ for $z_\ell \in Y - X \cap Y$, outputs the permutation of $(\{F_{k_2}(s_k)\}_{k \in [m]}, \{F_{k_2}(s_\ell)\}_{\ell \in [n-m]})$.

If s_i are function values, then the simulation is identical to hybrid 0, else it is identical to hybrid 1. □

4 Instantiation Based on the DDH Assumption

We construct a linear complexity PSU protocol from the DDH assumption in Figure 3.

Theorem 4.1. *The above PSU protocol is secure in the semi-honest model assuming H is a random oracle and the DDH assumption.*

Proof. We exhibit simulators $\text{Sim}_{\mathcal{R}}$ and $\text{Sim}_{\mathcal{S}}$ for simulating corrupt \mathcal{R} and \mathcal{S} respectively, and argue the indistinguishability of the produced transcript from the real execution. Let $|X \cap Y| = m$.

Corrupt sender: $\text{Sim}_{\mathcal{S}}$ simulates the view of corrupt \mathcal{S} , which consists of \mathcal{S} 's randomness, input, output and received messages.

Intuitively, the crux of the security proof is to simulate sender's view without the knowledge of Y , more precisely, the knowledge of $X \cap Y$. Looking ahead, we will use RO to program the hash outputs for elements in $Y - X \cap Y$ naively, and program the hash outputs for elements in X in a special way to ensure the associated messages sent from \mathcal{R} are pseudorandom.

We argue the output of $\text{Sim}_{\mathcal{S}}$ is indistinguishable from the real execution. For this, we formally show the simulation by proceeding the sequence of hybrid transcripts, where T_0 is the real view of \mathcal{S} , and T_2 is the output of $\text{Sim}_{\mathcal{S}}$.

Hybrid₀: $\text{Sim}_{\mathcal{S}}$ simulates with the knowledge of Y .

- $\text{Sim}_{\mathcal{S}}$ chooses the randomness for \mathcal{R} , i.e., picks $b \xleftarrow{R} \mathbb{Z}_p$.
- RO queries: picks $h_i \xleftarrow{R} \mathbb{G}$ and sets $H(z_i) := h_i$.

Parameters:

- Common input: $\mathbb{G} = \langle g \rangle$, hash function $H : \{0, 1\}^* \rightarrow \mathbb{G}$.
- Input of sender \mathcal{S} : $X = \{x_1, \dots, x_n\}$.
- Input of receiver \mathcal{R} : $Y = \{y_1, \dots, y_n\}$.

Protocol:

1. \mathcal{R} picks $b \xleftarrow{\mathbb{R}} \mathbb{Z}_p$, then sends $(H(y_1)^b, \dots, H(y_n)^b)$ to \mathcal{S} .
2. \mathcal{S} picks $a \xleftarrow{\mathbb{R}} \mathbb{Z}_p$, sends $(H(x_1)^a, \dots, H(x_n)^a)$ to \mathcal{R} ; he also computes $(H(y_1)^b)^a, \dots, (H(y_n)^b)^a$ and sends its permutation Γ to \mathcal{R} .
3. \mathcal{R} computes $(H(x_1)^a)^b, \dots, (H(x_n)^a)^b$, then sets $v_i = 0$ iff the value is not in Γ .
4. \mathcal{R} with select vector (v_1, \dots, v_n) and \mathcal{S} with input $\{(x_i, \perp)\}_{i \in [n]}$ engage in one-sided OT.
5. \mathcal{R} obtains $X - X \cap Y$, and thus obtains the union $X \cup Y$.

Figure 3: DH-PSU

- $\text{Sim}_{\mathcal{S}}$ outputs h_i^b for $y_i \in Y$.


Clearly, $\text{Sim}_{\mathcal{S}}$'s simulation is identical to the real view.



Hybrid₁: $\text{Sim}_{\mathcal{S}}$ still simulates with the knowledge of Y , but slightly change the simulation of random oracle

- RO queries: picks $d_i, e_i \xleftarrow{\mathbb{R}} \mathbb{Z}_p$, sets $H(z_i) := h_i = h^{d_i} \cdot g^{e_i}$. Here, h is a fixed random group element.
- $\text{Sim}_{\mathcal{S}}$ outputs h_i^b for $y_i \in Y$.

The modification in RO simulation does not alter the view.



Hybrid₂: $\text{Sim}_{\mathcal{S}}$ simulates without the knowledge of Y , and changes the simulation in the final input.

- $\text{Sim}_{\mathcal{S}}$ outputs random f_i for $i \in [n]$.

We argue that the view in hybrid 2 and hybrid 3 are computationally indistinguishable. More precisely, given $(g, g^\alpha, g^\beta, g^\gamma)$, a PPT adversary \mathcal{A} (with knowledge of X and Y) implicitly sets \mathcal{R} 's randomness $b := \alpha$.

- RO queries: sets $h := g^\beta$; picks $d_i, e_i \xleftarrow{\mathbb{R}} \mathbb{Z}_p$, sets $H(z_i) := h^{d_i} \cdot g^{e_i}$.
- \mathcal{A} outputs $g^{\gamma d_i + \alpha e_i}$ for $i \in [n]$.

If $(g, g^\alpha, g^\beta, g^\gamma)$ is a DDH tuple, then the simulation is identical to hybrid 1, else it is identical to hybrid 2.

Corrupt receiver: $\text{Sim}_{\mathcal{R}}$ simulates the view of corrupt \mathcal{R} , which consists of \mathcal{R} 's randomness, input, output and received messages.

We argue the output of $\text{Sim}_{\mathcal{R}}$ is indistinguishable from the real execution. For this, we formally show the simulation by proceeding the sequence of hybrid transcripts, where T_0 is the real view of \mathcal{R} , and T_2 is the output of $\text{Sim}_{\mathcal{R}}$.

Hybrid₀: $\text{Sim}_{\mathcal{R}}$ simulates with the knowledge of X .

- $\text{Sim}_{\mathcal{R}}$ chooses the randomness for \mathcal{S} , i.e., picks $a \xleftarrow{\mathcal{R}} \mathbb{Z}_p$.
- RO queries: picks $h_i \xleftarrow{\mathcal{R}} \mathbb{G}$ and sets $H(z_i) := h_i$.
- $\text{Sim}_{\mathcal{R}}$ computes and outputs $\{h_i^a\}_{x_i \in X}$; computes $\{(h_i^b)^a\}_{y_i \in Y}$, outputs its permutation.


Clearly, $\text{Sim}_{\mathcal{R}}$'s simulation is identical to the real view.



Hybrid₁: $\text{Sim}_{\mathcal{R}}$ still simulates with the knowledge of X , but slightly changes the simulation of random oracle

- RO queries: for $z_i \notin Y$, picks $r_i \xleftarrow{\mathcal{R}} \mathbb{Z}_p$ and sets $H(z_i) := h_i = g^{r_i}$; for $z_i \in Y$, picks $d_i, e_i \xleftarrow{\mathcal{R}} \mathbb{Z}_p$, sets $H(z_i) := h_i = h^{d_i} \cdot g^{e_i}$. Here, h is a fixed random group element.
- $\text{Sim}_{\mathcal{R}}$ computes and outputs $\{h_i^a\}_{x_i \in X}$; computes $\{(h_i^b)^a\}_{y_i \in Y}$, outputs its permutation.

The modification in RO simulation does not alter the view.



Hybrid₂: $\text{Sim}_{\mathcal{R}}$ simulates without the knowledge of X , and changes the simulation in the final input.

- $\text{Sim}_{\mathcal{R}}$ computes $f_j = h_j^a$ for $x_j \in X - X \cap Y$ (note that the information of such x_j can be inferred from \mathcal{R} 's output), picks random f_k for $k \in [m]$ (implicitly assigned for elements in $X \cap Y$), outputs (f_1, \dots, f_n) ; then picks random f_ℓ for $\ell \in [n - m]$ (implicitly assigned to elements in $Y - X \cap Y$), outputs the permutation of $(\{f_k^b\}_{k \in [m]}, \{f_\ell^b\}_{\ell \in [n-m]})$.

We argue that the view in hybrid 2 and hybrid 3 are computationally indistinguishable. More precisely, given $(g, g^\alpha, g^\beta, g^\gamma)$, a PPT adversary \mathcal{A} (with knowledge of X and Y) implicitly sets \mathcal{S} 's randomness $a := \alpha$, picks $b \xleftarrow{\mathcal{R}} \mathbb{Z}_p$.

- RO queries: sets $h := g^\beta$; for $z_i \notin Y$, picks $r_i \xleftarrow{\mathcal{R}} \mathbb{Z}_p$ and sets $H(z_i) := h_i = g^{r_i}$; for $z_i \in Y$, picks $d_i, e_i \xleftarrow{\mathcal{R}} \mathbb{Z}_p$, sets $H(z_i) := h_i = h^{d_i} \cdot g^{e_i}$.
- \mathcal{A} computes $f_j = (g^\alpha)^{r_i}$ for $x_j \in X - X \cap Y$, computes $f_k = g^{\gamma d_k + e_k}$ for $k \in [m]$, outputs (f_1, \dots, f_n) ; then prepares $f_\ell = g^{\gamma d_\ell + e_\ell}$ for $\ell \in [n - m]$, outputs the permutation of $(\{f_k^b\}_{k \in [m]}, \{f_\ell^b\}_{\ell \in [n-m]})$.

If $(g, g^\alpha, g^\beta, g^\gamma)$ is a DDH tuple, then the simulation is identical to hybrid 1, else it is identical to hybrid 2.

□

References

A Diffie-Hellman Based PSI

We recall the ideal PSI functionality in Figure 4.

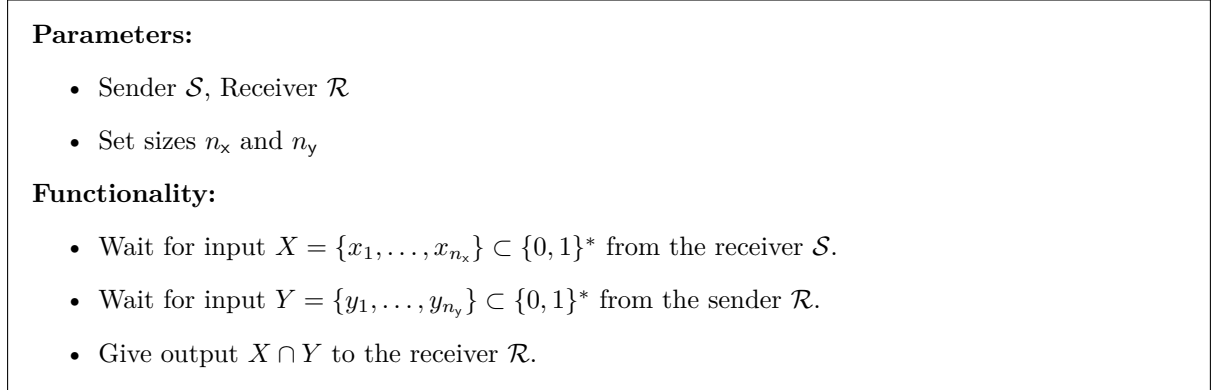


Figure 4: Private Set Intersection Functionality \mathcal{F}_{psi}

We first recall the DH-PSI in Figure 5.

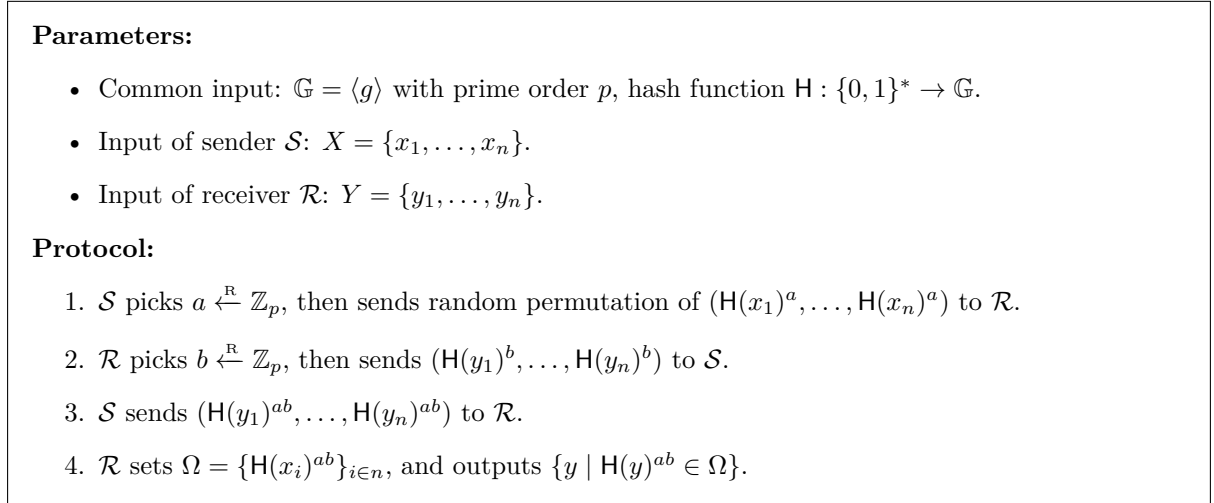


Figure 5: DH-PSI

Theorem A.1. *The above PSI protocol is secure in the semi-honest model assuming H is a random oracle and the DDH assumption.*

Proof. We exhibit simulators $\text{Sim}_{\mathcal{R}}$ and $\text{Sim}_{\mathcal{S}}$ for simulating corrupt \mathcal{R} and \mathcal{S} respectively, and argue the indistinguishability of the produced transcript from the real execution. Let $|X \cap Y| = m$.

Corrupt sender: $\text{Sim}_{\mathcal{S}}$ simulates the view of corrupt \mathcal{S} , which consists of \mathcal{S} 's randomness, input, output and received messages. $\text{Sim}_{\mathcal{S}}$ proceeds as follows.

We now argue the output of $\text{Sim}_{\mathcal{S}}$ is indistinguishable from the real execution. For this, we formally show the simulation by proceeding the sequence of hybrid transcripts, where T_0 is the real view of \mathcal{S} , and T_2 is the output of $\text{Sim}_{\mathcal{S}}$.

Hybrid₀: $\text{Sim}_{\mathcal{S}}$ simulates with the knowledge of Y .

- $\text{Sim}_{\mathcal{S}}$ chooses the randomness for \mathcal{R} , i.e., picks $b \xleftarrow{\mathbb{R}} \mathbb{Z}_p$

- RO queries: picks $h_i \xleftarrow{R} \mathbb{G}$ and sets $H(z_i) := h_i$.
- Sim_S outputs h_i^b for $y_i \in Y$ (message in step 2).


Clearly, Sim_S 's simulation is identical to the real view.



Hybrid₁: Sim_S still simulates with the knowledge of Y , but slightly change the simulation of random oracle

- RO queries: picks $d_i, e_i \xleftarrow{R} \mathbb{Z}_p$, sets $H(z_i) := h_i = h^{d_i} \cdot g^{e_i}$. Here, h is a random group element fixed at beginning.
- Sim_S outputs h_i^b for $y_i \in Y$.

The modification in RO simulation does not alter the view.



Hybrid₂: Sim_S simulates without the knowledge of Y , and changes the simulation in the final input.

- Sim_S outputs random f_i for $i \in [n]$.

We argue that the view in hybrid 2 and hybrid 3 are computationally indistinguishable based on the DDH assumption. More precisely, given $(g, g^\alpha, g^\beta, g^\gamma)$, a PPT adversary \mathcal{A} (with knowledge of X and Y) implicitly sets \mathcal{R} 's random tape $b := \alpha$:

- RO queries: sets $h = g^\beta$; picks $d_i, e_i \xleftarrow{R} \mathbb{Z}_p$, sets $H(z_i) := h_i = h^{d_i} \cdot g^{e_i}$.
- \mathcal{A} outputs $g^{\gamma d_i + \alpha e_i}$ for $i \in [n]$.

If $(g, g^\alpha, g^\beta, g^\gamma)$ is a DDH tuple, then the simulation is identical to hybrid 1 because $g^{\gamma d_i + \alpha e_i} = g^{\alpha \beta d_i + \alpha e_i} = h_i^b$, else it is identical to hybrid 2.

Remark A.1. From hybrid 1, we program $H(z_i) = h^{d_i} \cdot g^{e_i}$ to obtain a tight reduction by leveraging the random self-reducible property of the DDH assumption. Alternatively, we can simply program $H(z_i) = h^{d_i}$, then argue the indistinguishability using hybrid argument. As a result, the security reduction comes with a loose factor n .

Corrupt receiver: $\text{Sim}_{\mathcal{R}}$ simulates the view of corrupt \mathcal{R} , which consists of \mathcal{R} 's randomness, input, output and received messages. $\text{Sim}_{\mathcal{R}}$ proceeds as follows.

We now argue the output of $\text{Sim}_{\mathcal{R}}$ is indistinguishable from the real execution. For this, we formally show the simulation by proceeding the sequence of hybrid transcripts, where T_0 is the real view of \mathcal{R} , and T_2 is the output of $\text{Sim}_{\mathcal{R}}$.

The simulator has to emulate two messages:

- message in step 1: a permutation of $(H(x_1)^a, \dots, H(x_n)^a)$
- message in step 3: $(H(y_1)^{ab}, \dots, H(y_n)^{ab})$


Since the elements in $X - X \cap Y$ are unknown,

Hybrid₀: $\text{Sim}_{\mathcal{R}}$ simulates with the knowledge of X .

- $\text{Sim}_{\mathcal{R}}$ chooses the randomness for \mathcal{S} , i.e., picks $a \xleftarrow{R} \mathbb{Z}_p$
- RO queries: picks $h_i \xleftarrow{R} \mathbb{G}$ and sets $H(z_i) := h_i$.

- $\text{Sim}_{\mathcal{R}}$ outputs a random permutation of $(H(x_1)^a, \dots, H(x_n)^a)$ and $(H(y_1)^{ab}, \dots, H(y_n)^{ab})$.


Clearly, $\text{Sim}_{\mathcal{R}}$'s simulation is identical to the real view.



Hybrid₁: $\text{Sim}_{\mathcal{R}}$ still simulates with the knowledge of X , but slightly change the simulation of random oracle

- RO queries: for $z_i \in Y$, picks $r_i \xleftarrow{R} \mathbb{Z}_p$ and sets $H(z_i) := h_i = g^{r_i}$; for $z_i \notin Y$, picks $d_i, e_i \xleftarrow{R} \mathbb{Z}_p$, sets $H(z_i) := h_i = h^{d_i} \cdot g^{e_i}$. Here, h is a random group element fixed at beginning.

The modification in RO simulation does not alter the view.



Hybrid₂: $\text{Sim}_{\mathcal{R}}$ simulates without the knowledge of X , and changes the simulation in the output.

- $\text{Sim}_{\mathcal{R}}$ picks random f_j for $j \in [n - m]$ (associated with elements in $X - X \cap Y$), prepares $f_k = h_k^a$ for $k \in [m]$ (associated with elements in $X \cap Y$), then outputs a random permutation of $\{f_i\}_{i \in [n]}$, and $(H(y_1)^{ab}, \dots, H(y_n)^{ab})$.

We argue that the view in hybrid 2 and hybrid 3 are computationally indistinguishable based on the DDH assumption. More precisely, given $(g, g^\alpha, g^\beta, g^\gamma)$, a PPT adversary \mathcal{A} (with knowledge of X and Y) implicitly sets \mathcal{S} 's randomness $a := \alpha$, picks \mathcal{R} 's randomness $b \xleftarrow{R} \mathbb{Z}_p$:

- RO queries: sets $h = g^\beta$; for $z_i \in Y$, picks $r_i \xleftarrow{R} \mathbb{Z}_p$ and sets $H(z_i) := h_i = g^{r_i}$; for $z_i \notin Y$, picks $d_i, e_i \xleftarrow{R} \mathbb{Z}_p$, sets $H(z_i) := h_i = h^{d_i} \cdot g^{e_i}$.
- \mathcal{A} computes $f_j = g^{\gamma d_j + \alpha e_j}$ for $j \in [n - m]$, $f_k = (g^\alpha)^{r_k}$ for $k \in [m]$, then outputs a random permutation of $\{f_i\}_{i \in [n]}$ and $\{H(y_1)^{ab}, \dots, H(y_n)^{ab}\}$, where $H(y_i)^{ab} = (g^\alpha)^{r_i b}$.

If $(g, g^\alpha, g^\beta, g^\gamma)$ is a DDH tuple, then the simulation is identical to hybrid 1 because $g^{\gamma d_j + \alpha e_j} = g^{\alpha \beta d_i + \alpha e_i} = h_j^a$, else it is identical to hybrid 2.

□